Agentic Knowledge Graphs for Research Progression

PhD Qualifying Exam Presentation

Jason Cusati

Department of Computer Science
Virginia Tech
djjay@vt.edu

Spring 2026

2026-03-06 VT PhD Qualifier

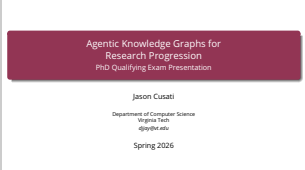

Good morning. I'm Jason Cusati, a PhD student working with Dr. Brown at the Code World No Blanket Lab here at Virginia Tech. The lab name is fitting, actually — today I'm arguing that AI research systems need somewhere to put their knowledge, because right now they're working without a net. Or a blanket. Today's talk is Agentic Knowledge Graphs for Research Progression — the academic way of saying: what if AI actually remembered what it learned, and could connect the dots — not just within a single paper, but across a whole field of research?

Outline

1. Motivation
2. Literature Survey
3. Synthesis
4. Proposal
5. Evaluation
6. Contributions
7. Conclusion

Here's the plan for the next 15 minutes. I'll start with the problem — why current AI-assisted workflows keep starting from scratch. Then I'll survey five domains and show where the gaps are. From there I'll present the proposal: a provenance-first architecture that gives agents a persistent, structured memory. I'll close with the evaluation framework and contributions. Questions at the end — I'm told there will be some, so I've been practicing my nervous smile.

The Structural Problem in AI-Assisted Research

- Research reasoning is **fragmented** and **transient**
- LLM-assisted workflows **amplify volatility**
- **Structural memory** is missing

The Structural Problem in AI-Assisted Research
- Research reasoning is **fragmented** and **transient**
- LLM-assisted workflows **amplify volatility**
- **Structural memory** is missing

Current AI systems dramatically accelerate document generation and research exploration.
However, they do not accumulate structured, validated, provenance-aware knowledge across iterations.
Outputs are transient. Each interaction effectively starts from scratch.
This work investigates how agentic systems combined with knowledge graphs can provide a persistent substrate for research progression.

The Structural Gaps

- **Persistence:** State beyond session boundaries
- **Traceability:** Linkage from Claim → Evidence
- **Evaluability:** Empirical measurement of progression

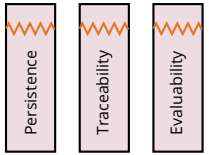

These are infrastructure problems, not model-quality problems.

The Structural Gaps
- **Persistence:** State beyond session boundaries
- **Traceability:** Linkage from Claim → Evidence
- **Evaluability:** Empirical measurement of progression
These are infrastructure problems, not model-quality problems.

This manifests as three specific structural gaps.
First, Persistence: We lack state that survives beyond a single chat session. Second, Traceability: When an agent makes a claim, we often cannot trace it back to the specific evidence that supported it. Third, Evaluability: We lack the tools to empirically measure if our research knowledge is actually progressing over time.
Crucially, these are infrastructure problems, not model quality problems. Better models won't fix this — better architecture will.

Guiding Question

How can research state be made persistent, traceable, and empirically evaluable in AI-assisted workflows?

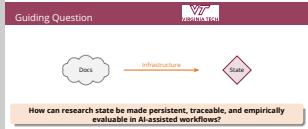

This leads to the guiding question of my proposal:
How can we transform this ephemeral interaction into something solid? How can research state be made persistent, traceable, and empirically evaluable in AI-assisted workflows?
To answer this, I've surveyed the current landscape to understand what pieces we already have — and what is missing.

Landscape Overview

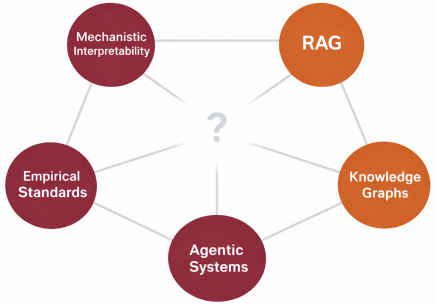

I surveyed five key domains that contribute to this problem.
We look at Mechanistic Interpretability to understand how models represent knowledge internally. We look at RAG and Knowledge Graphs as two ways of representing knowledge externally. We look at Agentic Systems which act on that knowledge. And finally, Empirical Standards to understand how we should evaluate these systems.

Mechanistic Interpretability

- Reverse-engineer the black box[a]
- Sparse features → human-interpretable structure
- Tells us *how* a model thinks

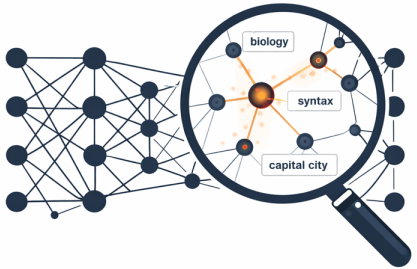

Limitation

[a] Sharkey et al. 2025.

Static weights only — no writable memory for new research state.

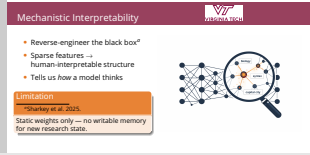

Starting with Mechanistic Interpretability, work by Sharkey et al. focuses on reverse-engineering the black box.
Techniques like sparse dictionary learning allow us to decompose model weights into understandable features. This is critical for understanding how a model thinks.
However, the limitation is that this focuses on the model's internal, static weights. It does not provide a mechanism for storing new research state that an agent generates during a workflow.

Concept-Based Interpretability

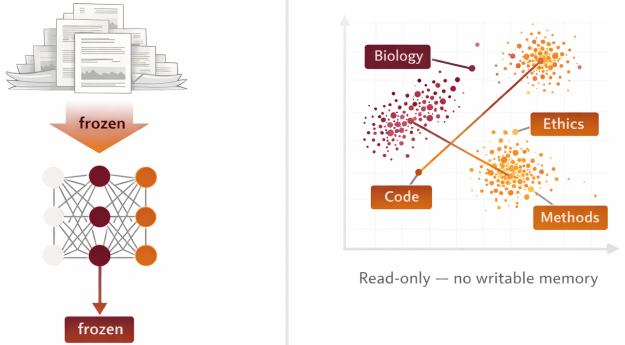

Limitation

Read-only — no mechanism to encode new research state.

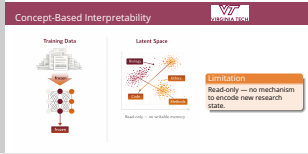

Concept-based interpretability takes a different angle. Instead of looking at circuits and weights, it asks: can we find directions in the model's latent space that align with human-readable concepts — like "Biology" or "Code"?

And the answer is kind of yes. Probes can find those directions. But what we're really measuring is what the model learned from its training data. It's sophisticated pattern matching on frozen knowledge.

The problem is that probing is fundamentally read-only. It tells us what the model baked in during training — not what we discovered in this session. Close the tab, and it's gone.

RAG & Vector Memory Systems

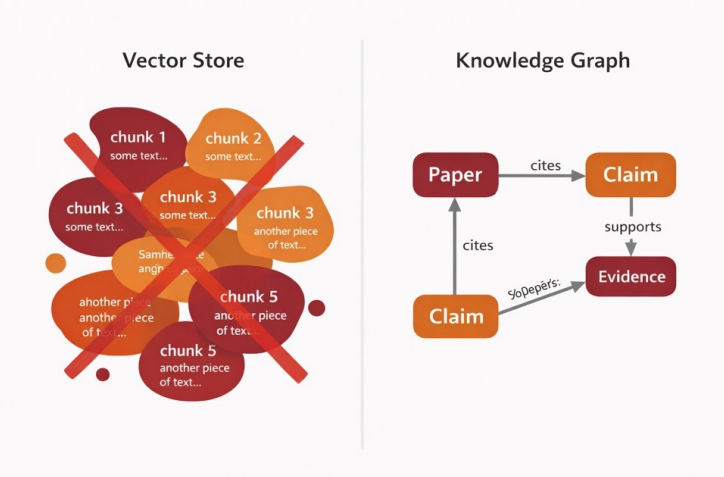

Limitations[a]

[a] Suryawanshi et al. 2025.

- No canonical memory
- No provenance chain
- No structured reasoning
- No longitudinal state

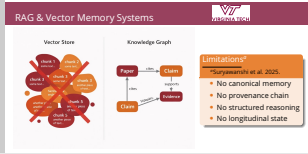

RAG is the go-to for external memory right now, and for good reason — it scales, it's fast, and semantic similarity gets you surprisingly far.
But Suryawanshi et al. point out the ceiling: vector retrieval gives you chunks, not structure. Similarity, not reasoning.
And critically, there's no canonical memory. Ask the same question after new papers drop and you get a different answer. No provenance chain — you can't trace a generated claim back to the specific text that justified it.

Knowledge Graph Research

Strengths

- Ontology-based structure
- Provenance modeling
- Logical relationships

Limitations

- Manual curation
- Static graphs
- No automated maintenance

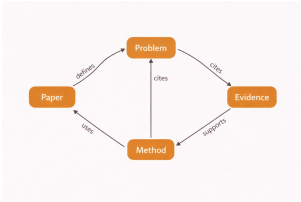

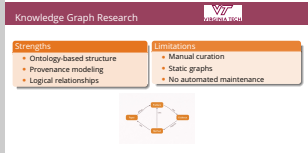

Knowledge Graphs offer exactly the structure that RAG lacks — explicit ontologies, logical relationships, and provenance modeling baked in. Suryawanshi et al. show how this enables cross-document reasoning that vector retrieval simply can't do.

But traditional KGs are high-friction. They require domain experts to curate them manually, they're static snapshots, and they don't evolve as the literature grows. A graph built today is already stale tomorrow.

The gap is automation. We have structure from KGs and we have automation from LLMs — but they're rarely integrated in a way that lets the graph actually evolve. That integration is exactly what this proposal targets.

Agentic Systems

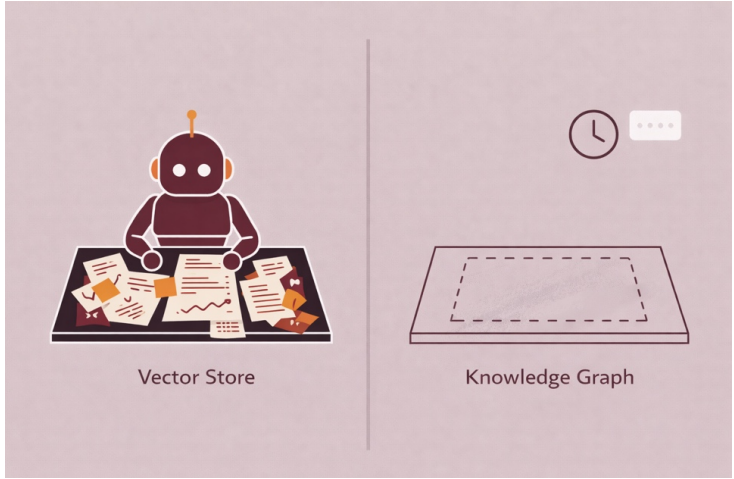

Limitations

- Ephemeral memory
- No structural persistence
- Weak auditability

This brings us to Agentic Systems — projects like Denario and the frameworks Trinkenreich et al. describe. These agents can execute complex, multi-step research tasks. That's genuinely impressive.

But they suffer from amnesia. Memory is bounded by the context window. Once the session ends, the structural state of the research is gone.

There's no persistent substrate anchoring their reasoning over time. Every session starts cold. That's the ceiling we need to break through.

Empirical Evaluation Standards

- Ralph et al. (2021): Empirical Standards for SE
- Baltes et al. (2025): Guidelines for LLM Studies

Observation

Single-pass or qualitative evaluations dominate.

Need

Longitudinal evaluation of coherent research state over time.

- [x] Standards defined
- [x] Guidelines published
- [ ] Agent validation?
- [ ] Longitudinal metrics?
- [ ] Provenance audit?

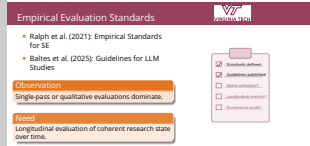

Finally, how do we evaluate any of this? Ralph et al. and Baltes et al. give us strong standards for SE and LLM experiments — clear, rigorous, and well-respected in the community.

But current agentic research frequently falls short. Evaluations are single-pass or qualitative. We test a snapshot, not a trajectory.

What's missing is longitudinal evaluation — frameworks that measure whether an agent is actually maintaining a coherent research state over weeks or months, not just getting the right answer once.

Cross-Cutting Gaps in the Literature

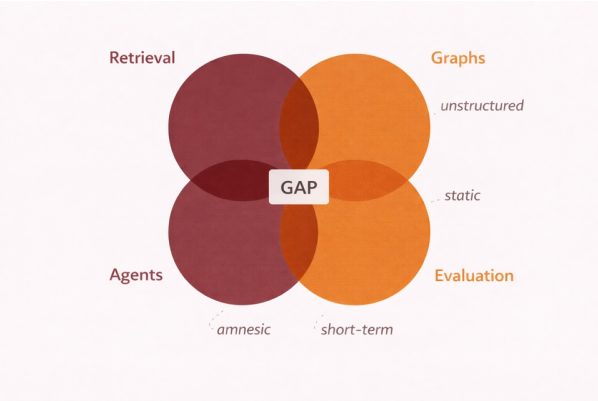

- **Retrieval** is unstructured
- **Graphs** are static
- **Agents** are amnesic
- **Evaluation** is short-term

This fragmentation motivates new infrastructure.

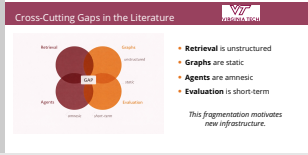

Synthesizing across all five domains, the fragmentation is clear.
Retrieval is unstructured. Graphs are static. Agents are amnesic. Evaluation is short-term. Each domain has made real progress in isolation — but none of them connect. Retrieval finds content but can't reason over it. Graphs provide structure but require human maintenance. Agents can execute complex tasks but forget everything when the session ends. And our evaluation tools weren't designed for longitudinal, multi-session workflows.
No existing system integrates persistent structured memory with autonomous agency and rigorous provenance. That gap in the center isn't an accident — it's the core blocker to true AI-assisted research progression.

Positioning My Work

VIRGINIA TECH

Intersection

Agentic systems + Knowledge graph persistence + Span-level traceability

Literature Refined:

- Stronger empirical evaluation
- SE domain scope
- Longitudinal validation

Positioning My Work
Intersection
Agentic systems + Knowledge graph persistence + Span-level traceability
Literature Refined:
Stronger empirical evaluation
SE domain scope
Longitudinal validation

My work positions itself directly in the gap we just mapped.
I'm proposing an architecture at the intersection of three areas: agentic systems, knowledge graph persistence, and span-level traceability. Prior work addresses these in isolation — RAG scales but lacks structure, knowledge graphs give structure but need manual curation, and agentic systems automate reasoning but suffer from amnesia. My proposal addresses all three simultaneously.
I also refine the evaluation approach from the literature. Instead of single-pass qualitative assessments, I enforce stronger empirical standards, narrow the scope to software engineering research tasks where the problem is most acute, and introduce explicit longitudinal testing to measure whether the system's knowledge actually persists over time.
This is not just combining existing tools — it is defining the infrastructure that makes their combination meaningful, reproducible, and auditable.

Proposed Architecture

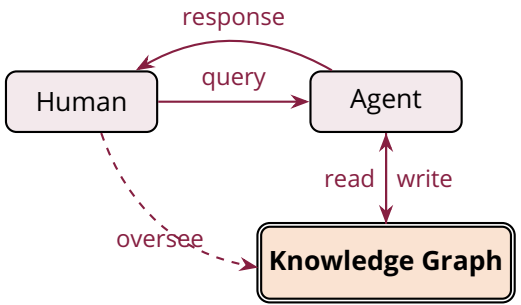

Persistent Research State Loop

- Human-in-the-loop oversight
- Agent-mediated updates
- KG as **Source of Truth**

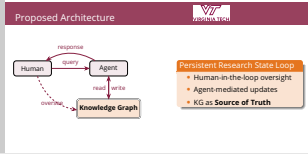

My proposed architecture is a “Persistent Research State Loop.”
Instead of a linear chat that disappears after the session, the human and agent collaborate to continuously update a shared, persistent Knowledge Graph.
The Graph is the source of truth. The agent reads from it to reason about the current state of a research problem, and writes to it to store new conclusions, evidence, and claims. Nothing gets lost between sessions.
The human stays in the loop — not just by reading final outputs, but by directly overseeing the state of the graph, correcting errors, and guiding the next query.
This is the architectural shift: from ephemeral conversation to persistent, structured, auditable research state.

Problem Extraction & Ranking

- Papers → problem nodes in KG
- Ranked by evidence density & citation weight
- Rank guides agent focus

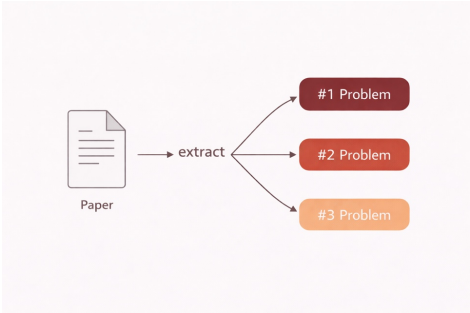

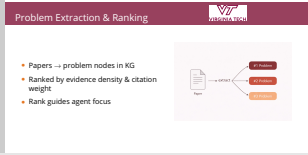

The agent doesn't just store information — it actively extracts research problems from ingested documents and represents them as first-class nodes in the Knowledge Graph. Those problems are then ranked by how much supporting evidence exists, how frequently they appear across sources, and their citation weight. This ranking is what guides the agent's next action — what to query, what to write, what to surface to the human.

It's a lightweight but critical step that turns passive retrieval into directed research progression.

Core Innovation

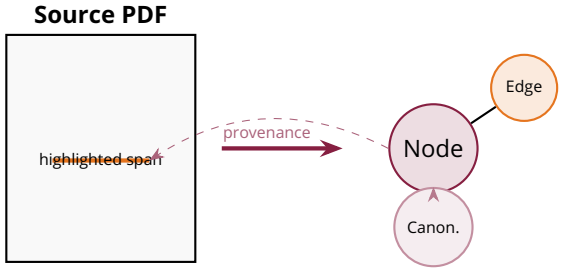

1. **Span-level alignment** — every node traces to a source sentence
2. **Canonicalization Agent** — de-duplicates entities in the graph
3. **Versioned provenance chains** — append-only audit trail

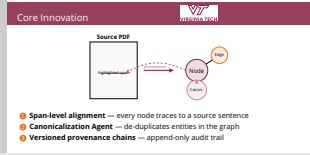

The core innovation is span-level evidence alignment.
Standard RAG retrieves vague chunks — a paragraph, a page, sometimes more. Our architecture links every node and edge in the Knowledge Graph back to the specific sentence or span in the source PDF. You always know exactly where a claim came from.
We also introduce a Canonicalization Agent whose sole job is to prevent fragmentation — ensuring we don't end up with five different nodes all representing "LLM." It merges duplicates and keeps the graph clean and queryable.
Finally, every write to the graph is versioned with a provenance chain — an append-only log recording the source span, the agent, and the prior graph state. Think of it like a git commit for research knowledge. That's what makes the system reproducible and defensible.

Evaluation Framework

Alignment Canonicalization Traceability Persistence

Claim	Metric	Baseline
Span-level alignment	Precision / Recall / F1	Standard RAG
Canonicalization	Cohen's κ	Embedding clustering
Traceability	Audit completeness %	RAG w/ citations
Persistence	Retention accuracy	Chat-only LLM

Key Question

Does the system retain knowledge better than a chat-only LLM across multiple sessions?

Evaluation Framework
Alignment Canonicalization Traceability Persistence
Claim Metric Baseline
Span-level alignment Precision / Recall / F1 Standard RAG
Canonicalization Cohen's κ Embedding clustering
Traceability Audit completeness % RAG w/ citations
Persistence Retention accuracy Chat-only LLM
Key Question
Does the system retain knowledge better than a chat-only LLM across multiple sessions?

To validate the system, I propose a four-claim evaluation matrix.
For span-level alignment, we measure Precision, Recall, and F1 against human-annotated span labels, with standard RAG as the baseline.
For canonicalization, we use Cohen's Kappa to measure agreement between the agent's entity merging and expert-labeled canonical clusters.
For traceability, we audit provenance completeness — what percentage of graph nodes have unbroken chains back to a source span.
And the headline metric: persistence. Does the system retain accurate knowledge across multiple sessions, compared to a chat-only LLM that resets each time? That's the claim no prior system has tested longitudinally.
These four metrics map directly to the four gaps from the literature.

Experimental Design Overview

- **Dataset:** Annotated SE spans
- **Injection:** Synthetic duplicates
- **Audit:** Traceability check
- **Task:** Retention measurement

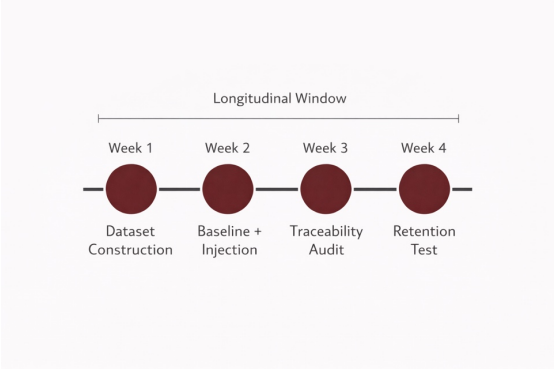

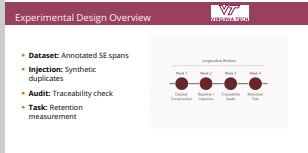

Important context: this evaluation happens *after* the system is built. The qualifier is a research proposal — this is what we would do to validate the system once implemented.
The experimental design is a 4-week longitudinal study.
Week 1: We construct the dataset — annotated SE papers with human-labeled spans as ground truth.
Week 2: We run the baseline and inject synthetic duplicates to stress-test the canonicalization agent — does it correctly merge entities a human would merge?
Week 3: Human traceability audit — are provenance chains intact as the graph grows?
Week 4: The retention test — we query the system for knowledge ingested in Week 1 and measure whether it retrieves it accurately. This is the longitudinal measurement no prior system has performed.
Four weeks is the minimum viable window for observing knowledge retention patterns. Ralph et al. and Baltes et al. both flag single-session evaluations as insufficient for systems claiming persistent behavior. The protocol naturally extends to longer periods if needed.

Contributions

Span-level Specification

Curation Protocol

Reference Infrastructure

1. **Span-level traceability specification** — foundational vocabulary for the field
2. **Agentic Curation Protocol** — how agents write, update, and validate the KG
3. **Reference infrastructure** — open-source implementation for benchmarking

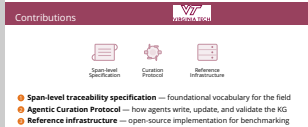

The expected contributions are threefold.
First, a span-level traceability specification — a precise technical definition of how every claim in the Knowledge Graph links back to its source span. This gives the field the vocabulary it needs to compare systems and measure provenance quality. The working definition of Research State itself is operationalized here: it is the versioned KG of claims, evidence links, and provenance chains.
Second, an Agentic Curation Protocol — a concrete specification for how agents should write to, update, and validate a shared knowledge graph. This is the operational layer that makes the architecture reproducible.
Third, a reference infrastructure implementation — open-source code and datasets that other researchers can build on, extend, or benchmark against.
Taken together, these move the field from vague aspiration to concrete, reusable artifacts.

Near-Term Work

- *Months 1–6:* Build system
- *Month 6–7:* Evaluation study
- *Month 7–8:* Artifact release
- *Year 1:* SE deployment

Near-Term Work
- *Months 1–6:* Build system
- *Month 6–7:* Evaluation study
- *Month 7–8:* Artifact release
- *Year 1:* SE deployment
Now Next
Build Evaluate Release Deploy

In the near term, my focus is on three sequenced steps.
Months 1 to 6: build the system — the KG pipeline, canonicalization agent, and provenance logging layer.
Month 6 to 7: run the controlled evaluation studies — the 4-week longitudinal study described on the previous slide.
Month 7 to 8: release all artifacts openly — the code, annotated datasets, and evaluation protocol — so others can reproduce and build on the work.
Year 1: SE-domain deployment — a working system applied to real software engineering literature review tasks, validating the proposal in a concrete research context.

Long-Term Vision

- Cross-domain expansion (Biology, Physics, Climate Science)
- Persistent AI research infrastructure
- **Cross-domain synthesis:** same problem, different fields

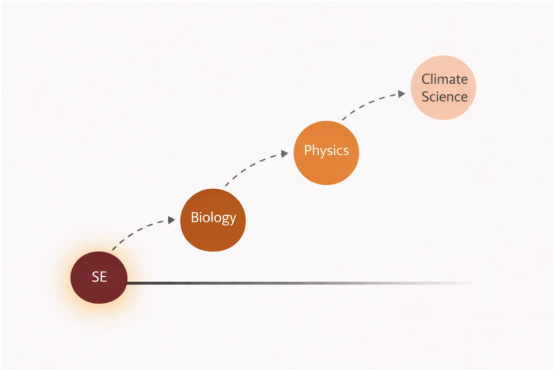

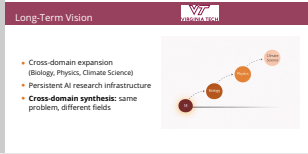

Long term, this vision extends far beyond Software Engineering.
Once the architecture is validated in the SE domain, the same infrastructure generalizes to any research discipline that accumulates claims across documents — Biology tracking experimental hypotheses, Physics linking theoretical claims to simulation evidence, Climate Science connecting model predictions to observational data.
But the more powerful long-term capability is cross-domain problem synthesis. Once knowledge graphs exist across multiple fields, the system can surface connections that communities would never find on their own. A reproducibility problem in SE may be structurally identical to a chain-of-custody problem in Chemistry, or a simulation provenance problem in Physics — different vocabulary, different venues, same underlying structure.
Infrastructure that doesn't just help one field think better, but helps science recognize when it's solving the same problem twice.

Conclusion

- **From:** Fragmented, ephemeral chat
- **To:** Persistent, traceable state
- **Goal:** Infrastructure for discovery

Infrastructure is the missing link to autonomous scientific discovery.

Conclusion
- **From:** Fragmented, ephemeral chat
- **To:** Persistent, traceable state
- **Goal:** Infrastructure for discovery

Infrastructure is the missing link to autonomous scientific discovery.

To conclude: we are moving from an era of fragmented, ephemeral chat interactions to an era of persistent, traceable research state — connected by Agentic Knowledge Graphs as the architectural bridge.

The core insight from this survey is that the gap in the literature is not a modeling gap. It is an infrastructure gap. Better models will not fix ephemeral memory. Better prompts will not give us span-level provenance. We need architecture that was explicitly designed to persist, trace, and evaluate research state across sessions and across time. Infrastructure is the missing link to AI-assisted discovery. That is what this proposal is building.

Thank you. I am happy to take your questions.

Questions?

Jason Cusati
djjay@vt.edu

Questions?
Jason Cusati
djjay@vt.edu

Backup Slides

Anticipated Questions — Motivation & Literature

Q [S2]: How is Evaluability different from accuracy? Accuracy measures a single output. Evaluability means the system exposes enough structure to measure *longitudinal* correctness — whether knowledge degrades, drifts, or accumulates errors across sessions.

Q [S7]: Have you considered hybrid RAG + KG approaches? Yes — hybrid systems exist (GraphRAG, etc.) but they use the KG only for retrieval routing, not as an auditable research state. Our contribution is provenance-first: every node is anchored to a source span regardless of retrieval method.

Q [S8]: How does your KG differ from existing domain KGs (Wikidata, biomedical)? Domain KGs are manually curated, static, and general-purpose. Ours is agent-maintained, append-only versioned, and scoped to the research workflow — claims link to specific paper spans, not to universal entity definitions.

Q [S9]: How does your system avoid the same amnesia you critique? By externalizing state into the KG. The agent's context window can reset; the graph does not. Every session reads from the graph first, so the agent starts informed, not

Anticipated Questions — Motivation & Literature

Q [S2]: How is Evaluability different from accuracy? Accuracy measures a single output. Evaluability means the system exposes enough structure to measure *longitudinal* correctness — whether knowledge degrades, drifts, or accumulates errors across sessions.

Q [S7]: Have you considered hybrid RAG + KG approaches? Yes — hybrid systems exist (GraphRAG, etc.) but they use the KG only for retrieval routing, not as an auditable research state. Our contribution is provenance-first: every node is anchored to a source span regardless of retrieval method.

Q [S8]: How does your KG differ from existing domain KGs (Wikidata, biomedical)? Domain KGs are manually curated, static, and general-purpose. Ours is agent-maintained, append-only versioned, and scoped to the research workflow — claims link to specific paper spans, not to universal entity definitions.

Q [S9]: How does your system avoid the same amnesia you critique? By externalizing state into the KG. The agent's context window can reset; the graph does not. Every session reads from the graph first, so the agent starts informed, not

Anticipated Questions — Architecture

Q [S13]: How does human oversight work in practice? The human reviews the graph state via a query interface — seeing which claims were added, what spans support them, and flagging incorrect merges. It is not continuous monitoring; it is checkpoint-based review.

Q [S13b]: How do you define a "research problem" vs. a finding or method? A problem node captures an open challenge stated or implied in the source. Findings are answers; methods are approaches. The extraction agent uses typed predicates to distinguish them in the ontology.

Q [S14]: How do provenance chains work? Each KG write records the source span, the agent, and the prior graph state — like a git commit log. An append-only versioned log; blockchain is a future scaling option.

Q [S14]: Why not just use RAG? RAG retrieves chunks, not structure. It has no canonical memory, no de-duplication, and no provenance. Our system adds all three.

Anticipated Questions — Architecture

Q [S13]: How does human oversight work in practice? The human reviews the graph state via a query interface — seeing which claims were added, what spans support them, and flagging incorrect merges. It is not continuous monitoring; it is checkpoint-based review.

Q [S13b]: How do you define a "research problem" vs. a finding or method? A problem node captures an open challenge stated or implied in the source. Findings are answers; methods are approaches. The extraction agent uses typed predicates to distinguish them in the ontology.

Q [S14]: How do provenance chains work? Each KG write records the source span, the agent, and the prior graph state — like a git commit log. An append-only versioned log; blockchain is a future scaling option.

Q [S14]: Why not just use RAG? RAG retrieves chunks, not structure. It has no canonical memory, no de-duplication, and no provenance. Our system adds all three.

Anticipated Questions — Evaluation

Q [S15]: How is your baseline RAG configured — does it also get citations? Yes — the RAG baseline includes citation-level retrieval (RAG w/ citations) so we are comparing against the strongest plausible alternative, not a strawman.

Q [S16]: Why 4 weeks? Why not longer? Four weeks is the minimum viable window to observe retention patterns; consistent with longitudinal SE field study norms (Ralph, Baltes). The protocol extends naturally to longer periods.

Q [S16]: What's your implementation timeline? Six months to build a working prototype, followed immediately by the 4-week evaluation study. Artifact release and SE deployment in Year 1.

Anticipated Questions — Evaluation

Q [S15]: How is your baseline RAG configured — does it also get citations? Yes — the RAG baseline includes citation-level retrieval (RAG w/ citations) so we are comparing against the strongest plausible alternative, not a strawman.

Q [S16]: Why 4 weeks? Why not longer? Four weeks is the minimum viable window to observe retention patterns; consistent with longitudinal SE field study norms (Ralph, Baltes). The protocol extends naturally to longer periods.

Q [S16]: What's your implementation timeline? Six months to build a working prototype, followed immediately by the 4-week evaluation study. Artifact release and SE deployment in Year 1.

Anticipated Questions — Contributions

Q [S17]: What does the formal Research State definition look like? It is operationalized through the architecture: Research State is the versioned KG of claims, evidence links, and provenance chains. Formalizing it as a published specification is Contribution 1.

Q [S18]: What is the biggest technical risk? Canonicalization quality. If the entity merging agent creates false merges, the graph loses distinguishing structure. The evaluation's Cohen's Kappa check directly stress-tests this.

Q [S19]: How do you handle terminology mismatches across domains? Cross-domain synthesis is a long-term goal requiring shared ontologies or embedding-based alignment. Near-term, each domain graph uses its own vocabulary; cross-domain linking is deferred to future work.

Anticipated Questions — Contributions

Q [S17]: What does the formal Research State definition look like? It is operationalized through the architecture: Research State is the versioned KG of claims, evidence links, and provenance chains. Formalizing it as a published specification is Contribution 1.

Q [S18]: What is the biggest technical risk? Canonicalization quality. If the entity merging agent creates false merges, the graph loses distinguishing structure. The evaluation's Cohen's Kappa check directly stress-tests this.

Q [S19]: How do you handle terminology mismatches across domains? Cross-domain synthesis is a long-term goal requiring shared ontologies or embedding-based alignment. Near-term, each domain graph uses its own vocabulary; cross-domain linking is deferred to future work.

References

References I

- Baltes, Sebastian et al. (2025). “Guidelines for Empirical Studies in Software Engineering involving Large Language Models”. In: *arXiv preprint arXiv:2508.15503*. See also llm-guidelines.org.
- Ralph, Paul et al. (2021). “Empirical Standards for Software Engineering Research”. In: *arXiv preprint arXiv:2010.03525*. Version 0.2.0. ACM SIGSOFT Paper and Peer Review Quality Task Force.
- Sharkey, Lee et al. (2025). “Open Problems in Mechanistic Interpretability”. In: *arXiv preprint arXiv:2501.16496*.
- Suryawanshi, Sudhanshu et al. (2025). “A Knowledge Graph-based RAG for Cross-Document Information Extraction”. In: *Proceedings of the 5th International Conference on Pervasive Computing and Social Networking (ICPCSN-2025)*. IEEE. DOI: 10.1109/ICPCSN65854.2025.11035854.
- Trinkenreich, Bianca et al. (2025). “Get on the Train or be Left on the Station: Using LLMs for Software Engineering Research”. In: *arXiv preprint arXiv:2506.12691*.

References I
Baltes, Sebastian et al. (2025). “Guidelines for Empirical Studies in Software Engineering involving Large Language Models”. In: arXiv preprint arXiv:2508.15503. See also llm-guidelines.org.
Ralph, Paul et al. (2021). “Empirical Standards for Software Engineering Research”. In: arXiv preprint arXiv:2010.03525. Version 0.2.0. ACM SIGSOFT Paper and Peer Review Quality Task Force.
Sharkey, Lee et al. (2025). “Open Problems in Mechanistic Interpretability”. In: arXiv preprint arXiv:2501.16496.
Suryawanshi, Sudhanshu et al. (2025). “A Knowledge Graph-based RAG for Cross-Document Information Extraction”. In: Proceedings of the 5th International Conference on Pervasive Computing and Social Networking (ICPCSN-2025). IEEE. DOI: 10.1109/ICPCSN65854.2025.11035854.
Trinkenreich, Bianca et al. (2025). “Get on the Train or be Left on the Station: Using LLMs for Software Engineering Research”. In: arXiv preprint arXiv:2506.12691.

References II

Villaescusa-Navarro, Francisco et al. (2025). “The Denario Project: Deep Knowledge AI Agents for Scientific Discovery”. In: *arXiv preprint arXiv:2510.26887*. Code: https://github.com/AstroPilot-AI/Denario.

References II

Villaescusa-Navarro, Francisco et al. (2025). “The Denario Project: Deep Knowledge AI Agents for Scientific Discovery”. In: *arXiv preprint arXiv:2510.26887*. Code: https://github.com/AstroPilot-AI/Denario.